

# PDC Project

Rumman Qadir 22i-1204

Awais Ahmed 22i-1225

Ahmed Asif 22i-1299

## 1. Problem Statement

We aim to explore efficient ways to enumerate all permutations of  $n$  elements build a graph connecting each permutation to its “parents” under  $n-1$  tree constructions, and record the full data set. The core challenges are:

- **Combinatorial explosion:**  $n!$  grows very quickly ( $10!=3.6$  million,  $11!=39.9$  million).
  - **Memory footprint:** Storing all permutations and parent links can exceed RAM.
  - **Compute throughput:** Even with on-the-fly generation, the total work is  $O(n! \times n)$
  - **Parallelization:** We must leverage multi-node (MPI) and multi-core (OpenMP) to finish in reasonable time.
- 

## 2. Solution Overview

We implemented several strategies:

### 1. Permutation Generation

- **Lexicographic (“factoradic”) indexing:** direct random access in  $O(n^2)$ .
- **In-place swapping** (Heap’s algorithm or adjacent swaps) to generate next permutation sequentially.

### 2. Parent Computation

- **logic:** to find “parents” by adjacent transpositions in  $O(n)$ .

- Encoded as `parent1(...)` and `FindPosition(...)` routines.

### 3. Index ↔ Permutation Conversion

- **Lehmer code** (factorial number system) for reversible index ↔ permutation.

### 4. Data Storage

- **In-RAM**: store all nodes and parent pointers in a contiguous table for fastest access.
- **On-SSD**: fall back to buffered text output when RAM is insufficient, writing in batches to avoid I/O stalls.

### 5. Parallelization

- **MPI**: split the full  $[0, n!)$  index space across ranks.
- **OpenMP**: within each rank, use `#pragma omp parallel for` on the hot loop to utilize all local cores.
- Ensured thread-safe operations (either via per-thread buffers or lock-free parent recomputation).

---

## 3. Methods Tried

Approach	Memory	Speed	Notes
<code>std::next_permutation</code> (sequential)	$O(1)$ gen	Slow	Simplicity, but single-threaded
Heap's Algorithm	$O(1)$ gen	Slow	Eliminates allocs, minor gain
Factoradic indexing + Lehmer decode	$O(n)$ per	Fast random	Best random-access, $O(n^2)$ per perm

Delta encoding (SJT deltas)	$O(n!)$	Compressed	4 bits per step, needed recompression for random access
RAM table + MPI + OpenMP	$O(n! \times n)$	<b>Fastest</b>	
SSD-streaming output	$O(1)$	I/O-bound	Useful if RAM < dataset


## 4. Implementation (RAM + MPI + OpenMP)

```
// Sketch of the hot loop (each rank works on [lo,hi]):
#pragma omp parallel for schedule(dynamic)
for (long long k = lo; k < hi; ++k) {
    auto perm = get_kth_permutation(k, n);    // Factoradic →  $O(n^2)$ 
    Record &R = table[k - lo];                // Pre-allocated array
    // copy node
    for (int i = 0; i < n; ++i) R.node[i] = perm[i];
    auto inv = inverse(perm);
    int r = pos(perm);
    // compute n-1 parents
    for (int t = 1; t < n; ++t) {
        auto p = parent1(perm, t, identity, inv, r);
        for (int i = 0; i < n; ++i)
            R.parents[t-1][i] = p[i];
    }
}
```

## 5. Performance Results

- **Serial MPI (1 slot)**
  - **Total permutations:** 39,916,800
  - **Time:** 297.418 s
  - **Throughput:** 134 k permutations/s

- **MPI only (4, 8, 4 slots)**
  - **Ranks:** 16, each ~2,494,800 perms
  - **Max time:** ~45.649 s
  - **Throughput:** 874 k permutations/s
- **OpenMP+MPI (2, 4, 2 slots)**
  - **Ranks:** 8, each ~4,989,600 perms
  - **Max time:** ~45.665 s
  - **Throughput:** 874 k permutations/s

 **~6.5× speed-up** vs serial by utilizing cluster + multicore.

---


## 6. Visual Comparison

Below is the **Performance Comparison** table and two charts showing

### 1. Total Execution Time



## 2. Throughput (records/s)



## 7. Conclusion & Lessons

- **In-RAM storage** with direct index mapping and parent recomputation yields the best throughput, at the cost of  $O(n! \times n)$  memory.
- **MPI + OpenMP** gives linear scalability up to the number of cores across machines.
- **SSD streaming** remains an option when RAM is insufficient, but becomes **I/O-bound**.
- Future work: GPU offload for the parent-finding logic and more advanced compression (delta + arithmetic coding) for extremely large  $n$ .